

POST-GROKING COLLAPSE AT THE REPRESENTATION–READOUT INTERFACE IN MUON-TRAINED TRANSFORMERS

Anonymous authors

Paper under double-blind review

ABSTRACT

Post-grokking collapse can destroy accuracy while preserving decodable task information. We study modular-arithmetic transformers trained with Muon on hidden weights and AdamW on embeddings and the output head. At five captured transitions, adjacent swaps localize the acute failure to the head update. Its mean-dependent logit term reproduces each failing test accuracy within 0.33 percentage points. Training-only linear decoders recover 98.20–100% at collapsed checkpoints. Corrected cross-entropy prevents joint train/test collapse through step 100,000 in five matched branches; an RMS-normalized model still fails through an embedding update. The captured updates locally descend accurate loss and overshoot at full scale. Finally, task-aligned Fourier filtering averages logits over inputs sharing the answer, allowing a memorizer to achieve perfect filtered accuracy.

1 INTRODUCTION

Generalization acquired during grokking can later collapse. We follow individual failures in modular-arithmetic transformers to identify the damaging update and determine what information survives. At a captured failure, dominant-frequency membership is unchanged and the task-family power distribution retains cosine similarity 0.9899 (Figure 1). Fresh decoders recover high accuracy from collapsed representations.

We use Muon (Jordan et al., 2024) with the routing used in large-scale training (Liu et al., 2025): hidden matrices receive orthogonalized momentum, while embeddings and the output head remain with AdamW. Prior work reports faster Muon grokking and stability tradeoffs (Tveit et al., 2025; Wang, 2026), connects inaccurate cross-entropy derivatives to feature inflation and amplified updates (Liu et al., 2026), and uses representation/readout decompositions and linear probes to study grokking (Chou et al., 2026).

Adjacent swaps isolate the readout update in the unnormalized model: the preceding head retains accuracy on updated features, and training-only decoders recover 98.20–100% at collapsed checkpoints. The mean-dependent head update nearly reproduces each failure. Matched freezes and corrected cross-entropy branches test prevention, while an RMS control identifies an embedding failure with accurate loss derivatives. The role of numerical error and the parameter group responsible for collapse differ between the two architectures.

Full-family Fourier filtering averages representations over operand pairs with the same answer. A lookup-table memorizer achieves perfect filtered accuracy despite chance-level held-out predictions. We therefore interpret filtering and rescaling as interventions on a task-informed projection. Raw-feature decoders and adjacent swaps establish surviving information independently of that projection.

Contributions.

- **Acute failure and surviving information.** Adjacent swaps isolate damaging readout updates in five unnormalized trajectories; training-only decoders recover high held-out accuracy, and a feature-mean intervention nearly reproduces each failure (§4).




Figure 1: A model and its filtered pathway failing, at ten-step evaluation resolution. **Left:** test accuracy falls from 100% to 19.04% and addition-filtered accuracy to 24.78%. The five dominant addition-frequency pairs are unchanged (Jaccard 1.0000), and power over all 56 addition pairs retains cosine similarity 0.9899. **Right:** accuracy returns to 99.90% within 300 steps, with power cosine 0.55. The same five pairs remain dominant, with a different allocation of power. The addition family consists of the non-constant (k, k) modes of the final residual’s operand-grid Fourier transform; accuracy filters also retain the constant mode.

Table 1: Parameter groups and routing. The AdamW baseline places all parameters under one AdamW instance. Selected settings: Muon hidden lr = 0.03, wd = 0.1, $\beta = 0.95$, five Newton–Schulz iterations; auxiliary lr = 10^{-3} , wd = 1.0; readout lr = 2.5×10^{-4} , wd = 1.0; AdamW baseline lr = 10^{-3} , wd = 3.0. Selection protocol and full sweeps in Appendix B.

Group	Parameters	Optimizer
Hidden	QKV, attention output, MLP input, MLP output, per block	Muon
Auxiliary	Token embedding, position embedding	AdamW
Readout	Unembedding	AdamW

- **Numerical mechanism and architectural boundary.** Matched arithmetic corrections test prevention over fixed horizons. Captured updates overshoot accurate local descent directions, and an accurate-loss RMS control fails through its embeddings (§4.4).
- **Scope of Fourier evidence.** An averaging identity and memorizer controls delimit the interpretation of full-family filtering and rescaling (§5, §6).

2 SETUP

Task. We train on $(a+b) \bmod p$, and on $(a-b) \bmod p$ where stated, using token sequences $[a, b, =]$ and p output classes. Defaults are $p = 113$ and a fixed random 30% training split of the p^2 operand grid. Training is full batch with cross-entropy. Original sweeps evaluate every 100 steps; diagnostic branches use finer grids where specified. The arithmetic follow-ups add the monitoring in §4.4.

Architecture. The model is a decoder-only transformer, $d_{\text{model}} = 128$, four heads, MLP width 512, ReLU, sequence length 3, every projection bias-free, and no normalization layers. An invertible change of residual basis can be absorbed by the matrices that write to and read from it. This is a symmetry of the parameterization; optimizer updates need not preserve it. The gain-free RMS control restricts this symmetry to $O(d)$. Depth is varied by appending blocks under nested initialization.

Routing and definitions. Parameters split into the three groups of Table 1; Muon receives the hidden group. Original grokking speeds use the first of five consecutive evaluations at or above 95% test accuracy. *Strictly stable* means no subsequent evaluation below 95%; we also report sub-threshold counts and minima. Matched branches fork from a single in-memory state with zero parameter

difference. CPU seed replications have identical logged metrics before the freeze (Appendix A). Follow-up confirmation uses six checks (Appendix J.1).

3 MUON GROKS FASTER, AND THE SOLUTIONS DO NOT HOLD

All nine Muon configurations in a 3×3 sweep of the hidden group reach sustained generalization, against seven of eleven AdamW configurations, with mean times of 13,011 and 30,486 steps; the mean is computed over AdamW’s successes alone, so the factor of 2.34 is conservative. On the selected settings the advantage is 1.54 at the median over five seeds, and Muon’s grokking step spans one hundred steps where AdamW’s spans seven thousand seven hundred. The ratios reach 4.8–4.9 under half width and a smaller modulus. AdamW never groks in the 20% training variant, giving a lower bound of 6.6; depth also increases the advantage (§7).

All nine Muon configurations lose test accuracy after grokking; six have a logged evaluation with both train and test accuracy below 95%. None is strictly stable: sub-threshold counts range from 2 to 592, post-grokking minima reach 0.78%, and the counts have no monotone relationship to either hidden hyperparameter. Four of the seven grokking AdamW configurations also collapse, with severity rising monotonically in the learning rate, from none at 10^{-3} to 918 sub-threshold evaluations at 10^{-2} . Sweeping the readout learning rate over a factor of five and its weight decay from zero to 0.5 changes the timing and depth of collapse and eliminates none of it; the zero-decay case removes the weight-decay pull on the readout and leaves the failure in place. Across five seeds neither selected setting is strictly stable: Muon records 137 to 321 sub-threshold evaluations on every seed with minima to 16.27%, and the AdamW reference records one or two on four of five, with a minimum of 27.59%. The tested Muon subtraction run also loses test accuracy; its matched AdamW run stays above threshold. Full sweeps, the baseline-selection protocol, and generality variants are in Appendix B.

4 LOCALIZING THE COLLAPSE AT THE INTERFACE

4.1 MUON UPDATES REMAIN NEARLY CONSTANT AT SMALL GRADIENT NORMS

Figure 5 measures steps 44,010–44,700 of an instrumented rerun. Training loss is near 1.5×10^{-7} , with hidden and non-hidden gradient norms of order 10^{-6} and 10^{-5} . Muon normalizes its momentum before a fixed Newton–Schulz iteration, reducing sensitivity to its magnitude. Across this window, the log–log slope of the applied hidden step against the hidden gradient norm is -0.03 . Slopes for the embedding and readout steps against the combined non-hidden norm are about $+1.5$, with R^2 across the three fits ranging from 0.84 to 0.98. The hidden step and decay largely cancel; their net displacement grows by 64% from step 44,100 to 44,700 while the gradient-driven step changes little. By step 44,700, root-mean-square displacement per parameter from the branch state is 0.073 for hidden parameters and 0.0091 for non-hidden parameters, a factor of 8.0. These measurements establish different displacement scales; their causal interpretation requires the numerical controls in §4.4. In the separate ten-step replay used for Figure 1 and Table 11, training loss rises by 3.8×10^8 across one interval, and training and test accuracy fall from 100% to 21.12% and 19.04%.

4.2 WHICH PARAMETERS THE FAILURE REQUIRES

Table 2 reports branches from a common state, each run for 2,000 steps with a different subset of parameters held fixed. Freezing either group suppresses the failure for the whole window; freezing the token or position embedding delays it without preventing it; freezing the unembedding prevents it over the same window, and leaving only the hidden matrices and the unembedding free reproduces it. With the non-hidden parameters frozen, Muon continues to apply hidden updates near 0.146 while loss remains near 1.5×10^{-7} .

These branches establish suppression over a finite window; freezing also changes subsequent gradients and optimizer history. To isolate the failing update, Figure 3a swaps features and heads between adjacent checkpoints at five captured events. In every seed, the old head retains or improves accuracy on the updated features, while the new head sharply lowers accuracy on the old features. Fresh linear decoders fitted using only training examples recover 98.20–100% at selected collapsed checkpoints

Table 2: Matched branches from a single state at step 44,000, each run to step 46,000. Collapse is recorded when training accuracy falls below 90%. Freezing either optimizer-managed side of the interface prevents it over this window; among non-hidden parameters only freezing the unembedding does, and hidden matrices together with the unembedding reproduce it.

Branch	First train collapse	Min test	Final test	Max non-hidden grad.
<i>Group branches</i>				
Control	44,700	16.77%	100.00%	913.07
Freeze hidden	none	99.07%	99.08%	2.79×10^{-5}
Freeze non-hidden	none	100.00%	100.00%	0
<i>Non-hidden component branches</i>				
Control	44,710	43.57%	99.93%	329.29
Freeze token embedding	44,910	34.74%	99.90%	3278.60
Freeze position embedding	44,760	37.78%	99.92%	323.76
Freeze unembedding	none	100.00%	100.00%	4.94×10^{-6}
Hidden and unembedding only	45,290	40.80%	100.00%	348.00




Figure 2: The intervention at five seeds. Before the freeze, the paired arms have identical logged metrics. One arm continues under ordinary Muon; the other holds embeddings and unembedding constant. The unfrozen arm falls below threshold 137 to 321 times per seed with minima as low as 16.27%; the frozen arm records none and never falls below 95.67%.

(Appendix J.2). Longer branches still show coadaptation: each decodes its own addition-filtered representation and fails on the other’s, and initially identical heads reach cosine similarity -0.028 . These observations establish readout sensitivity and surviving information. Exact global basis equivalence fails for the measured pairs, despite successful recovery of planted transformations. The joint normalization/orthogonalization ablation also has a separate limitation: non-finite loss was observed in all six generalizing runs (Appendix D).

4.3 FREEZING EMBEDDINGS AND READOUT STABILIZES THE TESTED DEPTH-1 RUNS

We hold the token embedding, the position embedding, and the unembedding constant from roughly two thousand steps after sustained generalization, while hidden matrices continue to receive Muon updates; we call the configuration *Stable Muon*. Across five runs spanning the main condition, $p = 97$, training fraction 20%, and width 64, a total of 451,400 post-grokking steps and 4,519 evaluations,




Figure 3: Follow-ups on five selected unnormalized trajectories. **(a)** Adjacent swaps: subscripts 0 and 1 denote states before and after the captured update. The old head retains accuracy on updated features; the new head sharply lowers accuracy on old features. **(b)** Training-only decoders recover 98.20–100% at collapsed snapshots; seeds 3 and 4 use later snapshots than (a). **(c)** Every original trajectory has a captured joint train/test failure below 90%; accurate CE from the matched step-6,000 state yields 0/5 through 100,000. Original failures were captured before step 30,000. Panels reuse seeds. Counts report binary incidence under mixed original monitoring and complete joint-event detection in corrected branches (Appendix J).

no evaluation falls below 95%, and in every run the post-grokking minimum is the accuracy at the evaluation that first met the criterion. Figure 2 pairs the arms at five seeds: freezing removes every sub-threshold evaluation on every seed, against an AdamW reference that achieves strict stability on one seed of five. Grokking speed is preserved because the freeze follows the threshold, and per-step cost falls afterward. In the long component comparison, holding only the unembedding fixed leaves eighteen of 936 evaluations below threshold; also freezing the embeddings removes them (Appendix E). Arithmetic correction also reaches the original 100,000-step budget (§4.4).

4.4 NUMERICAL CONTROLS AND THE FAILING STEP

Five matched branches replace stock cross-entropy with a cancellation-resistant `log1p` evaluation from step 6,000, after grokking and measurable derivative error. Through step 100,000, the original training horizon, none has a joint train/test failure below 90%; minimum measured test accuracy is 98.86%, and all finish at 100%. Every-update training checks trigger test evaluation below threshold, detecting every joint event. Brief test-only excursions can fall between grid evaluations. From step 15,000 to 20,000, accurate loss, target-derivative repair, and per-example zero-sum gradient projection each give 0/5 failures (Appendix J.3). Projection retains substantial derivative error while restoring the class-sum identity, connecting the result to numerical feature inflation (Liu et al., 2026).

The mean-dependent head update reproduces the acute failure. Before the captured failures, full-grid feature-mean norms reach 10,374–12,203. Let h_1 be the post-update residual, $W_{U,0}$ the preceding head, ΔW_U its actual displacement, and μ the training-set mean of h_1 . Starting from the rescued logits $W_{U,0}h_1$, the head-induced change decomposes exactly as

$$\Delta z = \Delta W_U(h_1 - \mu) + \Delta W_U\mu. \quad (1)$$

The same training-estimated μ is used for all test examples. Adding only $\Delta W_U\mu$ nearly reproduces each failure (Table 3); its test accuracy differs from the full update by at most 0.33 percentage points. The centered term retains 87.20–100% test accuracy and at least 99.09% training accuracy. This evaluation-only intervention establishes local sufficiency of the mean-dependent term, a class-

Table 3: Test accuracy (%) under the decomposition in Eq. 1. Each partial intervention adds one term to $W_{U,0}h_1$. The full update gives $W_{U,1}h_1$. All means use training examples only.

Seed	Preceding head	Mean term	Centered term	Full update
0	96.39	32.11	95.26	31.87
1	94.18	9.42	92.48	9.09
2	100.00	0.93	100.00	0.93
3	91.01	53.34	87.20	53.60
4	98.66	66.07	98.46	66.33

dependent logit offset shared by all examples. Feature growth and optimizer history remain coupled before the failing update.

Replacing gradients only at the terminal update still fails in all five seeds. Each complete update has negative directional derivative under the accurate training loss, and scales 0.001 and 0.01 lower that loss; the full update overshoots (Appendix J.4).

The corrected subtraction run also remains free of joint failure through 100,000 steps. A gain-free RMS pair bounds its scope: with accurate loss, test accuracy falls from 96.21% to 0.88%. Updating only embeddings gives 0.94%, while updating only the readout retains 96.21%. Corrected cross-entropy therefore leaves a distinct embedding failure possible; this one-seed counterexample does not estimate its prevalence (Appendix J.5).

5 TASK-ALIGNED FOURIER PROJECTIONS

We view each cached state as a function of the operand pair. An orthonormal two-dimensional Fourier transform over all p^2 pairs partitions its modes into six disjoint families: non-constant *addition* (k, k) , *subtraction* $(k, -k)$, *a-only*, *b-only*, the constant mode, and generic interaction. Family *sufficiency* retains that family’s modes; *ablation* removes them. Both interventions retain the constant mode when targeting a non-constant family. Unless stated otherwise, Fourier accuracies and solved-checkpoint selection ($\geq 95\%$) use the full operand grid. Appendix F specifies the interventions.

The filter supplies task structure. Retaining all (k, k) modes and the constant mode is exactly the orbit average

$$(P_{\text{add}}h)(a, b) = \frac{1}{p} \sum_{t=0}^{p-1} h(a+t, b-t), \quad (2)$$

with indices modulo p . Every donor has the same sum, and the linear readout commutes with this average. For subtraction, the shifts are $(a+t, b+t)$. A lookup-table predictor that outputs the correct one-hot vector on training examples and zero elsewhere becomes perfect after addition filtering whenever every answer orbit contains a training example. Appendix K verifies this control on the experimental splits. Full-family filtered accuracy thus measures the answer-conditioned average of the predictor; high accuracy alone leaves the learned algorithm unidentified.

Task-aligned projection preserves accuracy at the sampled solved checkpoints. In Table 4, the addition family reaches 100.00% at every selected addition checkpoint. Removing it leaves near-chance accuracy, while removing the subtraction family has little effect. Under matched initialization, configuration, and code path, training on $(a-b) \bmod 113$ exchanges these roles: $(k, -k)$ reaches 100% and (k, k) gives chance accuracy. These results characterize the task-aligned projection and family removal; both interventions retain the constant mode. The averaging identity also produces this operation dependence, so identifying shared learned algorithms requires additional evidence (Chughtai et al., 2023).

Frequency and phase matter within the family. Transporting the family onto other diagonal frequencies at matched power gives 2.64%. Randomizing phase per conjugate pair gives 0.19%, and randomizing phase per channel gives 0.78%. These interventions establish sensitivity to frequency assignment and phase. The memorizer controls in Appendix K also exhibit sparse filtered sufficiency and perturbation sensitivity, so these tests leave algorithm identification unresolved. Two branches

Table 4: Family sufficiency and ablation on the two operations. Addition columns at depths 2 and 4 aggregate eleven solved checkpoints; depth-1 columns are matched to the subtraction runs on initialization, configuration, and code path. Chance is $1/113 = 0.885\%$; across forty-three solved checkpoints spanning five seeds and three regimes, the task-aligned family gives exactly 100.00% at every one (Appendix F).

	Addition, depths 2 and 4	Addition, depth 1	Subtraction, depth 1
(k, k) sufficiency	100.00%	100.00%	0.88%
$(k, -k)$ sufficiency	0.88%	0.88%	100.00%
(k, k) ablation	1.34%	1.51%	99.68%
$(k, -k)$ ablation	98.26%	97.51%	2.54%
a -only / b -only / constant suff.	0.88%	0.88%	0.88%
Generic sufficiency	1.38%	1.59%	2.38%
Solved checkpoints	11	6	4




Figure 4: **Left:** task-aligned family accuracy (including the retained constant mode) against full-model accuracy over 98 checkpoints from the original studies, both evaluated through the current readout. *Projection failure* denotes low filtered accuracy; *projection masking* sits at the lower right, where the task-informed projection is accurate and the full model fails. The panel’s circuit labels refer to these operational outcomes. At 95 of 98 checkpoints the family alone is at least as accurate as the full model, and the three exceptions exceed it by under a point. **Right:** rescaling the non-constant task-aligned component of the masked checkpoint by α , with the rest of the representation and the readout untouched.

diverging from a bit-identical state each decode their own addition-filtered representation at 100% and the other’s at chance, showing dependence on the associated readout.

The update rule shapes the spread. On matched runs the Muon representation occupies 326 effective conjugate pairs of the non-constant spectrum against AdamW’s 5, leaving the task-aligned family holding 28% of non-constant power against AdamW’s 91%, and the contrast replicates across seeds and inside the family itself (Figure 6; Appendix G). These operand-frequency measurements complement prior work on Muon’s spectral properties (Ruan et al., 2026; Wang et al., 2025; Wang, 2026). Ablating the normalization and orthogonalization together collapses the count to 4, past AdamW’s, while task-family filtered accuracy remains 100%. These comparisons show a change in spectral dispersion. The filtered accuracy retains the averaging interpretation above.

6 INTERFERENCE IN THE TASK-ALIGNED PROJECTION

Filtering the final residual gives an answer-conditioned average through the current readout (Eq. 2). Comparing it with the full model measures interference from the remainder. In *projection failure*, filtering gives 2.65% and then chance, and the family has fallen out of the top twenty power pairs almost entirely. Fresh-decoder recovery in §4.2 provides a separate test of surviving task information. In *projection masking*, the filtered representation achieves 100% accuracy over the full operand grid.

Table 5: Correct-class margin at the masked checkpoint, decomposed by linearity of the bias-free readout. The competitor is the highest-scoring incorrect class under the full model, fixed across components. The constant mode belongs to the remainder. The final column is the fraction of examples with a positive component margin.

Branch	Component	Mean	On correct	On wrong	Fraction > 0
Control	full	-0.71	2.15	-3.13	0.4585
Control	task-aligned	+5.72	5.30	6.08	1.0000
Control	remainder	-6.43	-3.15	-9.21	0.0475
Freeze	full	+11.61	11.76	-1.67	0.9889
Freeze	task-aligned	+19.28	19.30	17.18	1.0000
Freeze	remainder	-7.67	-7.54	-18.86	0.0622

The full model reaches 45.85%. The projection averages to a correct answer for every orbit; the full prediction also includes the remainder. One matched trajectory passes through failure, full recovery, and masking; its frozen counterpart retains high filtered accuracy and substantial family power (Appendix H).

Masking is a competition on amplitude. Because the unembedding is linear and bias-free, the correct-class margin decomposes exactly into the task-aligned contribution and the remainder (Table 5). Against the full model’s highest-scoring incorrect class, the non-constant family contributes a positive margin on every example in both branches, including every example the masked model answers wrongly, and on those examples its margin is higher, +6.08 against +5.30. The failures coincide with stronger opposition from the remainder: the remainder contributes negatively on more than nine tenths of examples in both branches, and in the control an opposition of -6.43 outvotes a task-aligned term of +5.72. Rescaling the task-aligned component by α , with nothing retrained, raises the masked model monotonically on the sampled grid from 45.85% to at least 95% at $\alpha = 2.55$ and 99.91% at $\alpha = 6$: the intervention changes the balance of the two logit contributions without retraining. The rescaled component is $P_{\text{add}}h - \bar{h}$, with $\bar{h}$ the global mean. Recovery therefore depends on the task structure supplied by the projection. Family power share is a descriptive measure of this change.

Family power during grokking. Along an AdamW trajectory sampled every thousand steps, the addition family reaches 95.14% at step 3,000, when the model scores 0.30% and the family holds 1.05% of non-constant power. Ablating the family leaves near-chance accuracy at every sampled checkpoint. Its share rises to 70.62% at the test-accuracy jump, alongside greater concentration within the family. The memorizer control shows why this filtered accuracy alone cannot establish algorithm learning at step 3,000. Increasing power share describes the trajectory and is compatible with the cleanup account of Nanda et al. (2023). The margin decomposition and rescaling intervention establish amplitude competition between projected logit components at the masked checkpoint.

Spectral stability and functional change. Exact nonzero Fourier support is invariant under an invertible residual-basis map; power and its rankings are invariant under orthogonal maps. General invertible maps can change dominant-frequency membership. A jointly compensated basis change preserves predictions. At the collapse in Figure 1, dominant-frequency Jaccard is 1.0000 and power cosine is 0.9899, while the filtered pathway’s mean margin changes from +18.14 to -46.85. Over this interval the unembedding moves by 1.285 of its norm and the addition-filtered residual by 0.392; the old head on new features scores 26.31%, against 5.74% for the reverse. Adjacent swaps provide finer localization (Figure 3a). Accuracy returns to 99.90% within 300 steps with a changed power allocation. In the depth-4 pair, patching the frozen branch’s addition component of the final MLP write into the control gives 99.99% with the frozen head and 5.05% with the control head. These results show read-out sensitivity; exact basis equivalence requires the separate tests in Appendix J.2.

7 CONSTRUCTION AND ALIGNMENT SEPARATE IN DEPTH

The speed advantage is $2.22\times$ at depth 2 and exceeds $5.7\times$ at depth 4, where AdamW never sustains generalization within the budget. The AdamW configuration stable at depth 1 records 556

sub-threshold evaluations at depth 2. Averaging five phase perturbations at each stage, the first accuracy loss exceeding half the baseline occurs at an MLP activation in ten of eleven solved checkpoints and an MLP write in the remaining one. Earlier attention writes can be necessary under zeroing while insensitive to these frequency and phase interventions. This distinction limits localization by ablation alone. Across 41 checkpoints selected by full-grid accuracy, direct decoding of a post-block residual reaches 95% only at the final block and never exceeds 17.97% earlier. Thus compatibility with the final readout is concentrated at the last block in the sampled models. At depth 4, anchoring the interface cuts sub-90% evaluations about twenty-fold on an exact-state matched branch (Appendix I).

8 RELATED WORK

Grokking on modular arithmetic. Power et al. (2022) introduced the phenomenon. Nanda et al. (2023) reverse-engineered a Fourier algorithm, used causal ablations, and introduced restricted and excluded losses to track circuit formation and cleanup. Zhong et al. (2023) show that hyperparameters select qualitatively different algorithms; Chughtai et al. (2023) find shared algorithmic structure alongside variation across learned circuits.

Failure after grokking. Thilak et al. (2022) report cyclic instability under adaptive optimizers; Varma et al. (2023) describe ungrokking under reduced data; Prieto et al. (2025) identify softmax collapse and a loss-minimization direction that scales logits. Liu et al. (2026) connect finite-precision class-sum violations to coupled feature/classifier-mean growth, amplified Adam updates, and loss spikes; they test zero-sum corrections and normalization. Our arithmetic controls extend this account under Muon/AdamW routing. Adjacent swaps and the mean-dependent displacement intervention locate acute failures; fresh decoders measure surviving information. Joint train/test collapse differs from the perfect-training regime of anti-grokking (Prakash & Martin, 2025). The accurate-loss RMS control identifies an embedding failure in our normalized setting.

Muon. Jordan et al. (2024) introduced the optimizer and Liu et al. (2025) deployed it at scale under the routing used here. Tveit et al. (2025) and Wang (2026) report faster grokking in their tested settings; the latter studies orthogonalization and the speed/stability tradeoff with iteration count, including accuracy dips. We examine the five-iteration configuration over longer horizons. Ruan et al. (2026) and Wang et al. (2025) report higher effective rank and a more isotropic weight spectrum under Muon. Our Fourier measurements concern dispersion over operand frequencies, and the joint ablation tests the contribution of normalization and orthogonalization.

Representation and readout. Elhage et al. (2021) describe residual-stream basis freedom. Anthes et al. (2023) identify readout misalignment as a major component of forgetting under new tasks. Chou et al. (2026) use same-split linear probes to diagnose readout miscalibration during grokking and study relative representation and readout learning speeds, while Gomezjurado Gonzalez (2026) finds a decoder bottleneck in one-step Collatz prediction. Here the task stays fixed and the dissociation occurs after generalization: joint train/test failure can coexist with high accuracy under a fresh decoder.

9 DISCUSSION

The tests establish surviving task information, with training-only decoders reaching 98.20–100%, while rejecting exact global basis equivalence for the measured pairs. Arithmetic corrections suppress joint failures over fixed horizons; captured updates overshoot accurate local descent directions. Together with adjacent swaps and the feature-mean intervention, these results support a numerical contribution to acute readout failure. The relative contributions of feature growth and optimizer history remain unresolved.

The accurate-loss RMS embedding failure limits this account. The main controls cover five selected trajectories; the operation and normalization extensions use one seed each. Behavior beyond 100,000 steps and brief test-only excursions remain unresolved, and AdamW collapses limit attribution specifically to Muon. Full-family Fourier filtering supplies answer-orbit structure and mixes training and test inputs. The memorizer control rules out using perfect filtered accuracy alone as evidence of an arithmetic algorithm.

REPRODUCIBILITY STATEMENT

The original experiments run on a single device from released code; their repository (anonymized for review) contains the training and analysis scripts, 191 CSV run logs, and 20 HTML summaries. Follow-up artifacts add raw diagnostics, model/optimizer/random-state checkpoints, replay checks, and the scripts for every new table and figure. Matched freezes and arithmetic branches verify shared states before intervention. Appendix A specifies the original setup; Appendix J gives the decoder, basis, arithmetic, and architecture-control protocols and their fixed horizons.

AI USE STATEMENT

Generative AI tools assisted with the design and implementation of follow-up experiments, numerical verification, mathematical derivations, literature comparison, figure generation, and manuscript revision. The released protocols, code, checkpoints, and verification records document these analyses.

REFERENCES

- Daniel Anthes, Sushrut Thorat, Peter König, and Tim C. Kietzmann. Diagnosing catastrophe: Large parts of accuracy loss in continual learning can be accounted for by readout misalignment. In *Conference on Cognitive Computational Neuroscience*, pp. 748–751, 2023. URL <https://arxiv.org/abs/2310.05644>.
- Chi-Ning Chou, Oscar Uzdelewicz, Neng-Chun Chiu, Yao-Yuan Yang, and SueYeon Chung. Two speeds of learning: A representation-readout decomposition of grokking and double descent. *arXiv preprint arXiv:2605.27078*, 2026.
- Bilal Chughtai, Lawrence Chan, and Neel Nanda. A toy model of universality: Reverse engineering how networks learn group operations. In *International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2302.03025>.
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. URL <https://transformer-circuits.pub/2021/framework/index.html>.
- Laura Gomezjurado Gonzalez. The long delay to arithmetic generalization: When learned representations outrun behavior. *arXiv preprint arXiv:2604.13082*, 2026.
- Keller Jordan, Yuchen Jin, Vlado Boza, Jiacheng You, Franz Cesista, Laker Newhouse, and Jeremy Bernstein. Muon: An optimizer for hidden layers in neural networks, 2024. <https://github.com/KellerJordan/Muon>.
- Hanqing Liu, Jianjun Cao, Yuanze Li, and Zijian Zhou. Grokking or glitching? how low-precision drives slingshot loss spikes. *arXiv preprint arXiv:2605.06152*, 2026. URL <https://arxiv.org/abs/2605.06152>.
- Jingyuan Liu et al. Muon is scalable for LLM training. *arXiv preprint arXiv:2502.16982*, 2025. URL <https://arxiv.org/abs/2502.16982>.
- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2301.05217>.
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*, 2022. URL <https://arxiv.org/abs/2201.02177>.
- Hari K. Prakash and Charles H. Martin. Grokking and generalization collapse: Insights from HTSR theory. In *HiLD Workshop at the International Conference on Machine Learning (ICML)*, 2025. URL <https://arxiv.org/abs/2506.04434>.
- Lucas Prieto, Melih Barsbey, Pedro A. M. Mediano, and Tolga Birdal. Grokking at the edge of numerical stability. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://arxiv.org/abs/2501.04697>.
- Tianyu Ruan, Fengzhuo Zhang, Shuche Wang, and Shihua Zhang. Muon learns more robust and transferable features than adam. *arXiv preprint arXiv:2606.09658*, 2026. URL <https://arxiv.org/abs/2606.09658>.
- Vimal Thilak, Etai Littwin, Shuangfei Zhai, Omid Saremi, Roni Paiss, and Joshua Susskind. The slingshot mechanism: An empirical study of adaptive optimizers and the grokking phenomenon. *arXiv preprint arXiv:2206.04817*, 2022. URL <https://arxiv.org/abs/2206.04817>.
- Amund Tveit, Bjørn Remseth, and Arve Skogvold. Muon optimizer accelerates grokking. *arXiv preprint arXiv:2504.16041*, 2025. URL <https://arxiv.org/abs/2504.16041>.

Vikrant Varma, Rohin Shah, Zachary Kenton, János Kramár, and Ramana Kumar. Explaining grokking through circuit efficiency. *arXiv preprint arXiv:2309.02390*, 2023. URL <https://arxiv.org/abs/2309.02390>.

Shuche Wang, Fengzhuo Zhang, Jiayang Li, Cunxiao Du, Chao Du, Tianyu Pang, Zhuoran Yang, Mingyi Hong, and Vincent Y. F. Tan. Muon outperforms adam in tail-end associative memory learning. *arXiv preprint arXiv:2509.26030*, 2025. URL <https://arxiv.org/abs/2509.26030>.

Yufeng Wang. The active ingredient in muon’s grokking. *arXiv preprint arXiv:2607.20512*, 2026. URL <https://arxiv.org/abs/2607.20512>.

Ziqian Zhong, Ziming Liu, Max Tegmark, and Jacob Andreas. The clock and the pizza: Two stories in mechanistic explanation of neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2306.17844>.

Table 6: Operational definitions. Evaluations occur every 100 steps, so a sustained criterion requires the threshold to hold across a 400-step span.

Term	Definition
Sustained 99.9% train	First of five consecutive evaluations at or above 99.9% train accuracy
Sustained 95% test	First of five consecutive evaluations at or above 95% test accuracy
Memorization plateau	Sustained-95 test step minus sustained-99.9 train step
Strictly stable	No post-grokking evaluation below 95% test accuracy

Table 7: Selected configurations. Stable Muon and Muon are the same configuration up to the freeze; their pre-freeze trajectories can differ through backend nondeterminism, which Section 4.3 reports.

Configuration	Group	Optimizer	Settings
AdamW baseline	All parameters	AdamW	$\text{lr} = 10^{-3}$, $\text{wd} = 3.0$
Muon	Hidden	Muon	$\text{lr} = 0.03$, $\text{wd} = 0.1$, $\beta = 0.95$, $\text{NS} = 5$
	Auxiliary	AdamW	$\text{lr} = 10^{-3}$, $\text{wd} = 1.0$
	Readout	AdamW	$\text{lr} = 2.5 \times 10^{-4}$, $\text{wd} = 1.0$
Stable Muon	Muon, with auxiliary and readout groups frozen after sustained generalization		

A SETUP DETAILS AND OPERATIONAL DEFINITIONS

Each example is the token sequence $[a, b, =]$; the vocabulary contains $p + 1 = 114$ tokens and the output layer $p = 113$ classes. The ordered operand grid contains $p^2 = 12,769$ examples; a fixed random 30% (unless stated) forms the training set, generated once from the run seed and shared by all matched comparisons. The original sweeps load saved initial states, evaluate every 100 steps, and checkpoint every 1,000. Their budgets are 100,000 steps at depths 1 and 2 and 300,000 at depth 4. Diagnostic branches use the shorter horizons and finer grids specified below.

Each block holds four two-dimensional matrices (fused QKV 384×128 , attention output 128×128 , MLP input 512×128 , MLP output 128×512), 196,608 parameters; with embeddings and unembedding the depth-1 model has 226,048 parameters. Depth is varied by appending blocks under nested initialization, so deeper models begin from a strict superset of the shallower model’s initial state. Muon maintains an exponential moving average of the gradient with $\beta = 0.95$, forms a Nesterov-style update, divides it by its Frobenius norm plus 10^{-7} , applies five quintic Newton–Schulz iterations with coefficients $(3.4445, -4.7750, 2.0315)$, scales by $\sqrt{\max(1, \text{rows}/\text{cols})}$, and applies decoupled weight decay; the implementation refuses any parameter that is not a matrix. All AdamW groups use $\beta_1 = 0.9$, $\beta_2 = 0.999$.

Table 6 gives the operational definitions for the original experiments; evaluations occur every 100 steps, so a sustained criterion holds across a 400-step span. Training sweeps and matched-branch experiments run at a single seed on an accelerator backend that is not bitwise deterministic across processes; the five-seed replications run on CPU, where execution is bitwise deterministic, and the Muon and Stable Muon arms have identical logged metrics over the seventy-four to seventy-five evaluations preceding the freeze. This check compares recorded metrics; it does not compare every parameter tensor at each evaluation. The seed sets the initialization and the train/test split together, so the five seeds are five draws of the problem instance.

B SWEEPS, BASELINE SELECTION, AND GENERALITY

Tables 8 and 9 give both sweeps under identical architecture, split, and budget. The fastest AdamW configuration ($\text{lr} = 10^{-2}$, $\text{wd} = 1.0$) is unusable: 918 sub-threshold evaluations at this seed, and across seeds three of five never reach sustained 95%, the two that do record 805 and 956, and all five end near chance. The baseline is therefore the fastest configuration that is strictly stable in this sweep, $\text{lr} = 10^{-3}$, $\text{wd} = 3.0$; the seed replication shows strict stability at that setting is a property of one seed rather than of the configuration. For Muon no configuration is strictly stable, so the selected configuration is the fastest.

Table 8: AdamW sweep. Seven of eleven retained configurations reach sustained 95% test accuracy. Post-grokking evaluations below 95% increase monotonically with the learning rate.

LR	WD	Sustained 95%	Min after grokking	Evals < 95%	Strictly stable
3×10^{-4}	0.1	no grok			
3×10^{-4}	1.0	no grok			
3×10^{-4}	3.0	25,500	95.67%	0	yes
10^{-3}	0.1	no grok			
10^{-3}	1.0	52,500	96.49%	0	yes
10^{-3}	3.0	8,200	95.01%	0	yes
3×10^{-3}	0.1	49,500	13.26%	2	no
3×10^{-3}	1.0	8,500	53.72%	4	no
10^{-2}	0.1	62,900	0.54%	258	no
10^{-2}	1.0	6,300	0.41%	918	no
10^{-2}	3.0	no grok			

Table 9: Muon sweep over the hidden group. All nine configurations reach sustained 95% test accuracy and none is strictly stable. The count of sub-threshold evaluations has no monotone relationship to either hyperparameter.

LR	WD	Sustained 95%	Min after grokking	Evals < 95%	Strictly stable
0.03	0.1	5,400	69.40%	56	no
0.01	0.1	6,000	53.24%	107	no
0.03	0.3	7,200	75.63%	2	no
0.01	0.3	8,500	76.60%	145	no
0.03	0.03	9,000	0.78%	592	no
0.003	0.1	14,400	90.07%	5	no
0.003	0.03	19,000	21.24%	130	no
0.01	0.03	19,500	0.83%	473	no
0.003	0.3	28,100	78.59%	265	no

Both optimizers fit the training set between steps 200 and 800 in every grokking configuration, so the difference lies almost entirely in the delay from memorization to generalization: 30,100 steps for AdamW against 12,544 for Muon averaged over successes. Across five seeds Muon reaches sustained generalization between 5,300 and 5,400 steps and the AdamW reference between 4,000 and 11,700, medians 5,400 and 8,300.

Varying the readout learning rate over 10^{-4} , 2.5×10^{-4} , and 5×10^{-4} gives post-grokking minima of 37.13%, 79.40%, and 3.81%; varying its weight decay over 0, 0.1, and 0.5 gives 33.34%, 0.79%, and 1.12%; every setting is unstable. On $(a - b) \bmod 113$ the Muon run records 18 post-grokking evaluations below 95% with a minimum of 1.52% and the AdamW run records none, reproducing the depth-1 addition asymmetry at 194 and none. Table 10 varies modulus, training fraction, and width; Stable Muon and Muon share a configuration up to the freeze, so differences in their grokking steps bound run-to-run variation on this backend.

C THE QUIET WINDOW

Figure 5 uses the every-step instrumented rerun, with a quiet window from 44,010 through 44,700 and failure at 44,702. Table 11 uses a separate ten-step spectral replay, whose first observed failure is at 44,710. The following step/decay diagnostics use the instrumented rerun. On the hidden group the gradient-driven component is 0.1499 at step 44,100 against a decay component of 0.1475, a ratio of 0.984, leaving a net displacement of 0.0462 that grows to 0.0759 across the window while the orthogonalized step stays flat; over the window the hidden gradient norm grows by 3.92 while its gradient-driven step changes by 0.97, and the non-hidden gradient grows by 2.60 while the embedding and readout steps grow by 2.65 and 2.84. The readout decay component has larger norm than its gradient-driven component on 14.9% of quiet-window steps. In the ten-step spectral replay, the combined non-hidden gradient norm rises from 3.52×10^{-5} to 836.10; the orthogonalized

Table 10: Controlled variation of modulus, training fraction, and width. Stable Muon and Muon share a configuration up to the freeze, so differences in their grokking step measure pre-intervention run-to-run variation on this backend.

Variant	Regime	Sustained 95%	Min post-grokking	Outcome
$p = 97$, train 30%, width 128	AdamW	31,000	95.52%	stable
	Muon	6,300	15.49%	unstable
	Stable Muon	6,300	95.05%	stable
$p = 113$, train 20%, width 128	AdamW	no grok		
	Muon	15,100	49.25%	unstable
	Stable Muon	14,500	95.16%	stable
$p = 113$, train 30%, width 64	AdamW	84,800	96.38%	stable
	Muon	17,700	53.23%	unstable
	Stable Muon	16,500	95.67%	stable




Figure 5: Every-step measurements over 691 updates before a collapse. **Top:** loss and gradient norms in the quiet window. **Middle:** gradient-driven applied steps and log-log slopes against the hidden gradient norm for Muon and the combined non-hidden norm for both AdamW groups. **Bottom:** hidden step and decay magnitudes differ by under 2% at step 44,100. Their vector sum is smaller and grows by 64% from step 44,100 to 44,700 while the gradient-driven step changes little.

Muon step over that interval is 0.1429, lower than the 0.1455 preceding it, and its maximum over the branch occurs after the event.

D ABLATING THE NORMALIZED-ORTHOGONALIZATION PATH

The ablation removes Frobenius normalization and the Newton–Schulz iteration together, retaining momentum, shape scaling, and weight decay. In the tested sweep, learning requires a rate of 0.3 or 1.0. Six ablated runs generalize, with no recorded post-grokking evaluation below 95% (Table 12). All six were observed to terminate with non-finite loss, in float32 and float64, at steps 22,889–38,955. The matched Muon comparator records thirteen sub-threshold evaluations through step 22,800. Over its full 100,000 steps, its minimum and final accuracies are 29.49% and 96.48%. At the analyzed checkpoints, ablation reduces dispersion from 326.09 effective conjugate pairs to 4.11, while the task-aligned family retains 100% sufficiency (Appendix G).

E PREVENTION DETAILS

In four of the five runs the freeze is triggered automatically, 2,000 steps after the first evaluation of a five-evaluation streak at or above 95%; the main-condition run uses a fixed step. The first logged frozen evaluation lands 2,100 to 2,200 steps after sustained generalization; the actual freezes occur

Table 11: Diagnostics through the quiet window and across the collapse. The training loss and both gradient norms are numerically negligible until the failure. The Muon update column is the orthogonalized step before weight decay, varying by under three percent across the pre-collapse rows. These rows come from the ten-step spectral replay; the net-step diagnostics in the text use the separate instrumented rerun.

Step	Train loss	Hidden grad. norm	Non-hidden grad. norm	Muon update norm
44,100	1.77×10^{-7}	2.33×10^{-7}	1.95×10^{-5}	0.1498
44,300	1.55×10^{-7}	2.73×10^{-7}	2.20×10^{-5}	0.1494
44,500	1.50×10^{-7}	4.06×10^{-7}	2.67×10^{-5}	0.1481
44,700	1.53×10^{-7}	8.64×10^{-7}	3.52×10^{-5}	0.1455
44,710	57.50	25.26	836.10	0.1429
44,720	6.58	5.09	92.55	0.1655

Table 12: Ablated runs that reach sustained generalization. Minimum is the lowest test accuracy from the sustained-95 step to the last recorded evaluation. Terminal is the step at which non-finite loss was observed during the original run.

Learning rate	Backend	Precision	Grokked	Minimum	Terminal
0.3	MPS	float32	3,100	97.28%	23,514
0.3	CPU	float32	3,100	97.52%	24,485
0.3	CPU	float64	3,100	97.06%	22,889
1.0	MPS	float32	6,800	97.11%	23,800
1.0	CPU	float64	7,700	95.96%	29,831
1.0	CPU	float32	6,500	95.60%	38,955

100 steps earlier. Within a depth-2 run the median step costs 57.87 ms before the freeze and 54.72 after; at depth 4, 111.39 and 107.75; pre-freeze rates match ordinary Muon’s. Holding only the unembedding constant from the same step leaves the run unstable, minimum 83.13% and eighteen of 936 evaluations below 95%; freezing the embeddings and unembedding together at the same step produces none.

F FOURIER DECOMPOSITION AND INTERVENTION SUITE

For a chosen point in the forward pass the model runs on all p^2 operand pairs, the states form a tensor $X \in \mathbb{R}^{p \times p \times d}$, and an orthonormal two-dimensional DFT over the operand axes gives $\hat{X}[k, l, :]$ with power $\|\hat{X}[k, l, :]\|^2$. Because X is real, modes occur in conjugate pairs and every filter retains or removes both members together; transforms run in single precision on CPU. The families of Table 14 are pairwise disjoint and exhaust the grid, and addition and subtraction are disjoint because p is odd. Power fractions are normalized against non-constant power; every non-constant-family sufficiency filter also retains the constant mode, which becomes a fixed logit offset under the bias-free unembedding. Ablation removes only the selected non-constant modes, retaining this offset. The two masks are complementary on the non-constant spectrum and share the constant mode. In the margin decomposition and α sweep, the task-aligned component excludes the constant mode, which belongs to the remainder. Solved-checkpoint selection and intervention accuracies use the full operand grid; training trajectories explicitly labeled test accuracy use the held-out split.

Table 16 applies sufficiency and ablation to each family at forty-three solved checkpoints spanning five seeds and the Muon, Stable Muon, and AdamW regimes. Equal-power relocation ranges 0.88% to 15.04% over 430 replicates. Among the tested descending-power prefixes, the smallest reaching 95% contains two to four conjugate pairs at AdamW checkpoints (median 3), against three to seventeen at Muon (median 12.5) and eight to nineteen at Stable Muon (median 11). Along the AdamW trajectory of §6, the five strongest pairs give 4.91% at step 3,000, 19.48% at 20,000, 70.73% at 30,000, and 100% at 37,000, while the family’s share rises from 1.05% through 1.66% and 11.08% to 70.62%, reaching 90.51% by step 100,000; test accuracy over those steps is 0.30%, 0.29%, 2.08%, and 100%.

Table 13: Stable Muon across five runs covering four settings; the main condition appears twice, under a fixed freeze step and under the automatic trigger. The minimum is the lowest test accuracy at any evaluation from the sustained-95 step onward, and in every run it is the value at that step itself.

Run	Grokked	First frozen eval.	Post-grokking steps	Minimum	Evals < 95%
Main condition	5,900	8,100	94,100	95.34%	0
Selected Stable Muon	5,400	7,500	94,600	95.74%	0
$p = 97$	6,300	8,400	93,700	95.05%	0
Training fraction 20%	14,500	16,600	85,500	95.16%	0
Width 64	16,500	18,600	83,500	95.67%	0
Total			451,400		0

Table 14: The six mode families. The families are pairwise disjoint and exhaust the $p \times p$ grid, so any statement about one family and its complement is an exact accounting of the whole representation.

Family	Modes	Count at $p = 113$
Addition	$(k, k), k \neq 0$	112
Subtraction	$(k, -k), k \neq 0$	112
a -only	$(k, 0), k \neq 0$	112
b -only	$(0, l), l \neq 0$	112
Constant	$(0, 0)$	1
Generic interaction	all remaining (k, l)	12,320
Total		12,769

The full-family accuracies above include the orbit average in Eq. 2. Prefix filtering and phase or frequency perturbations measure additional spectral structure, but their interpretation as evidence of a learned algorithm requires controls against memorization. Appendix K supplies that comparison.

G SPECTRAL DISPERSION

Table 17 reports the entropy effective count $\exp(-\sum_i q_i \log q_i)$ over all 6,384 non-constant conjugate pairs and the task-aligned family’s power share on matched depth-1 runs. Here q_i denotes normalized pair power. Across the five-seed study AdamW occupies 4.40 effective pairs on average over thirty solved checkpoints, Muon 385.85 over twenty-five, and Stable Muon 343.96 over thirty-six: freezing the interface leaves the dispersion where Muon put it and removes the instability anyway. Within each family’s $(p-1)/2$ pairs, the generality analysis instead uses the inverse participation ratio $1/\sum_i q_i^2$, with q_i renormalized within the family: 1.99–2.59 for solved AdamW variants and 21.33–55.34 for Muon. On the uniform operand grid, a predictor using only the opposite family and a constant offset has accuracy exactly $1/p$, since sum and difference are independent for odd p . Across solved matched checkpoints, AdamW places 0.0–7.8% of non-constant power there and Muon 2.5–17.7%.

H THE MATCHED BRANCH THROUGH FAILURE, RECOVERY, AND MASKING

Table 18 follows the family’s top-twenty membership, power share, isolated accuracy, and full-model accuracy along the depth-4 matched pair of \$6 from a bit-identical state. Over 174,000 steps, the frozen branch retains high filtered accuracy and substantial family power; the control loses all three at filtered-pathway failure through its current readout, recovers all three, and loses the share again at projection masking while the isolated family still gives 100%. The α -sweep crossing at $\alpha = 2.55$ occurs at a share of 0.6247, below the frozen branch’s native 0.9069. Replaying a depth-1 collapse and recovery at ten-step resolution: full accuracy returns to 99.90% within 300 steps while the addition-family power cosine to the pre-collapse state reaches only 0.55; the pre-collapse representation holds between 4.2% and 5.7% under the recovered readout throughout, while the recovered representation reads at 90.73% under the pre-collapse readout. The matched branches’

Table 15: Interventions applied to a cached state. Preserved and altered quantities specify the intervention; an accuracy change measures sensitivity to that perturbation. Non-constant-family sufficiency and ablation both retain the constant mode.

Intervention	Preserved	Destroyed
Family sufficiency	Selected family and constant mode	Other non-constant modes
Family ablation	All other families	Modes in one family
Equal-power relocation	Total power, family membership	Which frequencies carry it
Random native subset	Pair count, native coefficients	Which of the used frequencies survive
Pair-global phase scramble	Support and power exactly	Relative phase across pairs
Channelwise phase scramble	Support and power exactly	Relative phase across channels
Cross-readout substitution	Both representations intact	Pairing of state with its own readout

Table 16: Family interventions at forty-three solved checkpoints, spanning five seeds and the Muon, Stable Muon, and AdamW regimes. Sufficiency retains the named family and constant mode; ablation removes the named non-constant family and retains the rest. Accuracies and selection use the full grid. The chance rate is 0.88%.

Family	Sufficiency	Range	Ablation	Range
Addition	100.00%	[100.00, 100.00]	1.71%	[0.67, 4.24]
Subtraction	0.88%	[0.88, 0.88]	99.61%	[95.79, 100.00]
<i>a</i> -only	0.88%	[0.88, 0.88]	99.85%	[97.76, 100.00]
<i>b</i> -only	0.88%	[0.88, 0.88]	99.85%	[97.72, 100.00]
Constant	0.88%	[0.88, 0.88]		
Generic interaction	3.28%	[0.78, 14.21]	99.99%	[99.78, 100.00]

five strongest modes are identical at the branch point and share at most one pair at sampled steps 160,000 through 230,000. Patching the addition component of the frozen branch’s final MLP write into the control under the frozen readout gives 99.99%; under the control’s own readout, 5.05%; patching its full final MLP write gives 97.61% and 4.46% under the two readouts; patching only the non-addition component of that MLP write into the control gives 36.79%, below the control’s unmodified 45.85%.

I DEPTH

Muon reaches sustained generalization in 2,300 steps at depth 2 against 5,100, and in 52,600 at depth 4 where AdamW peaks at 96.69% and never sustains the threshold within 300,000 steps. Table 19 compares steps and elapsed time; a Muon step costs 1.75–1.80× an AdamW step. At the depth-4 checkpoint in Table 20, zeroing the block-2 and block-3 attention writes leaves 0.94% and 0.44% accuracy. Frequency relocation and phase scrambling at these writes preserve near-baseline accuracy. Both perturbations cause a large drop at the block-3 MLP activation. Figure 7 and Table 20 show the reported perturbation realization. Later-stage outcomes vary across the five perturbations, so near-chance accuracy in this realization does not establish uniform sensitivity across perturbations.

Across eleven solved checkpoints, the first stage where mean pair-phase scrambling costs more than half the baseline accuracy is an MLP activation in ten and an MLP write in one. This criterion averages five perturbations per stage. Direct decoding reaches 95% only at the final block across all 41 checkpoints selected by full-grid accuracy, never exceeding 17.97% earlier (Table 21). Within the eleven-checkpoint stage-intervention subset, among seventeen residual states with direct-readout accuracy between 2% and 95%, family filtering improves sixteen, with a mean increase of 4.88 percentage points across all seventeen. Across eighteen depth-4 checkpoints after the shared state, between-branch power cosine stays above 0.98 in the first block; in the fourth it lies between 0.03 and 0.44 at fifteen checkpoints. The exact-state branch at step 126,200 raises mean post-branch accuracy from 48.07% to 97.50%, final accuracy from 36.88% to 98.87%, and cuts evaluations below 90% from 1,353 to 66.




Figure 6: The twenty highest-power conjugate pairs of the final residual at converged checkpoints, on the mode grid, sized by share of non-constant power; grey lines mark $l = k$ and $l = -k$. AdamW places 80.1% of non-constant power on a single pair of the family the task selects; Muon’s strongest pair carries 7.9%, with the rest spread across the family and beyond it; ablating the normalized path returns the strongest pair to 58.9%. Between the addition and subtraction panels the occupied line moves from one diagonal to the other.

Table 17: Spectral dispersion at solved checkpoints of the matched depth-1 runs. Effective pairs is $\exp(-\sum_i q_i \log q_i)$, with normalized pair powers q_i , over all 6,384 non-constant conjugate pairs, reported as a mean with range. Family power is the task-aligned family’s share of non-constant power, meaning (k, k) for addition and $(k, -k)$ for subtraction, and top-five power is the share carried by the five strongest pairs of the spectrum. The final row ablates the normalization and orthogonalization together.

Operation	Optimizer	Effective pairs	Family power	Top-five power
Addition	AdamW	4.95 [3.68, 10.20]	91.0%	93.7%
Addition	Muon	326.09 [96.91, 578.90]	28.0%	13.1%
Subtraction	AdamW	2.54 [2.45, 2.72]	95.4%	96.5%
Subtraction	Muon	211.46 [33.25, 377.24]	27.7%	16.0%
Addition	Muon, path ablated	4.11 [2.04, 6.05]	66.5%	94.0%

Table 18: The addition family along matched branches. Pairs counts how many of the twenty highest-power conjugate pairs at the final residual belong to the family; share is the family’s fraction of non-constant power; addition-only is accuracy when the representation is filtered to the family.

Step	Branch	Pairs	Share	Addition-only	Full
126,200	control	11	83.95%	100.00%	94.07%
126,200	freeze	11	83.95%	100.00%	94.07%
160,000	control	1	11.88%	2.65%	3.29%
160,000	freeze	12	89.33%	100.00%	98.97%
180,000	control	0	0.32%	0.88%	0.97%
180,000	freeze	12	88.32%	100.00%	97.95%
230,000	control	4	67.69%	100.00%	98.41%
230,000	freeze	12	90.78%	100.00%	99.47%
300,000	control	3	20.38%	100.00%	45.85%
300,000	freeze	12	90.69%	100.00%	98.89%

Table 19: Time to sustained generalization, in optimizer steps and in elapsed seconds on a single device. Per-step cost is the median over evaluation intervals. The AdamW baseline at depth 4 is right-censored at the step budget.

Depth	Optimizer	ms/step	Steps to grok	Seconds to grok	Advantage
2	AdamW	33.06	5,100	170	
	Muon	57.81	2,300	133	2.22× steps, 1.28× time
4	AdamW	61.68	> 300,000	> 19,768	
	Muon	110.87	52,600	5,856	> 5.70× steps, > 3.37× time

Table 20: Stage-level interventions at depth 4, step 53,000, baseline 99.30%. Relocation transports the addition family onto other diagonal frequencies at matched power; phase scrambling randomizes the phase of each conjugate pair. The two stochastic columns show the reported realization (replicate 4 of five); zeroing is deterministic.

Stage	Kind	Relocation	Phase	Zeroed
Block 1 attention write	write	99.30%	99.30%	86.31%
Block 2 attention write	write	99.33%	99.29%	0.94%
Block 3 attention write	write	99.29%	99.30%	0.44%
Block 3 MLP activation	activation	1.66%	0.00%	1.19%
Block 3 MLP write	write	0.95%	0.00%	1.19%
Block 4 attention write	write	3.58%	0.00%	58.44%

Table 21: Direct-readout accuracy by post-block residual, decoding with the unembedding and skipping the remaining blocks. Values are ranges over checkpoints at which the model has at least 95% full-grid accuracy.

Depth	Regime	Earlier blocks	Final block
2	AdamW	0.7–1.1%	100.00%
2	Muon	0.8–5.1%	95.8–100.00%
2	Stable Muon	1.3–5.2%	96.4–99.4%
4	Muon	0.8–18.0%	95.1–99.3%
4	Stable Muon	0.7–18.1%	95.2–98.7%




Figure 7: Depth-4 Muon at step 53,000. **Left:** the reported stage-intervention realization. Relocation and phase scrambling first cause a large drop at the block-3 MLP activation. The two stochastic curves use replicate 4 of five and illustrate sensitivity; individual later-stage outcomes vary across perturbations. Zeroing the earlier block-2 and block-3 attention writes also destroys accuracy, despite their low sensitivity to these Fourier perturbations. **Right:** direct decoding through the final unembedding, using each post-block residual or its addition-filtered component. Both reach high accuracy only at the final block.

J DECODER, BASIS, AND ARITHMETIC FOLLOW-UPS

The decoder and matched-branch tests use the original architecture, data splits, and optimizer settings; operation and normalization controls are specified separately below. The basis study adds 52,698 baseline updates and 9,000 continuation updates. The subsequent causal study adds 364,302 updates: five 24,000-step arithmetic branches, fifteen 5,000-step specificity branches, four original-trajectory extensions totaling 49,302 steps, and four 30,000-step operation or architecture controls. A subsequent horizon extension adds 420,000 updates: each of the five corrected addition branches and the corrected subtraction run continues from step 30,000 to 100,000. Verification replays are excluded from these totals. Multiple checkpoints and interventions from one trajectory are dependent observations.

J.1 MONITORING AND MATCHED STATES

The causal branches and fresh baseline replications use a joint event: training and test accuracy must both fall below 90% after sustained grokking. Training accuracy is checked at every update; test accuracy is checked every 100 updates and at every post-grokking state with training accuracy below 90%. Thus every joint event is detected, while brief test-only excursions can fall between evaluations. Six consecutive grid evaluations at or above 95% confirm grokking. The confirmation step is the *last* evaluation of this six-check streak. Original sweep speeds instead report the start of a five-check streak (§2); the criteria therefore differ as well as their reporting points.

Within each seed, matched arithmetic branches restore identical parameters, optimizer buffers, and random-number-generator state. Exact next-update replays reproduce all three. The original prefixes combine historical evaluation grids with densely captured failures; their positive events establish binary incidence. We therefore compare event presence, without comparing their counts, durations, or historical minima against the uniformly monitored branches. The five trajectories were selected for mechanistic study and give no estimate of collapse prevalence over an unrestricted seed population.

Table 22: Test accuracy (%) at adjacent captured failures. h_0, W_0 precede the update; h_1, W_1 follow it. Each head is evaluated with each feature state, without retraining. Original events are captured under mixed historical monitoring grids (Appendix J.1).

Seed	Before: $W_0 h_0$	Old head: $W_0 h_1$	New head: $W_1 h_0$	After: $W_1 h_1$
0	95.49	96.39	31.86	31.87
1	94.14	94.18	13.56	9.09
2	99.99	100.00	1.01	0.93
3	89.54	91.01	50.80	53.60
4	91.34	98.66	66.75	66.33

J.2 FRESH DECODERS AND THE BASIS HYPOTHESIS

Fresh bias-free linear decoders are fitted to the raw final residual using only the original 3,830 training examples. An uncentered SVD preconditioner retains all 128 dimensions, with singular values floored at 10^{-6} of the largest. Each inner fit estimates its own preconditioner. A 20% split of the training set selects an L_2 penalty on the preconditioned coefficients from $\{0, 10^{-8}, 10^{-6}, 10^{-4}, 10^{-2}\}$ by validation cross-entropy; the selected decoder is then refitted on all training examples. The 8,939 test examples enter only the final evaluation. Fits use float64 logistic regression and L-BFGS-B; convergence and unregularized controls are recorded.

At the selected collapsed checkpoints, decoder accuracies for seeds 0–4 are 98.72%, 98.20%, 100.00%, 98.77%, and 100.00%, against 99.51%, 99.56%, 100.00%, 99.40%, and 100.00% at their references. Probe checkpoints are steps 17,494, 17,721, 18,069, 16,309, and 16,060: the last two occur after the adjacent transitions in Table 22. These results establish substantial surviving linear task information; the small gaps in three seeds leave room for reduced decodability or probe limitations.

For column residuals, an exact global basis account requires one invertible R satisfying $h_1(x) = R h_0(x)$ across the operand grid. With examples stacked as rows, full column rank makes this equivalent to equality of the two feature column spaces. We test train-fitted linear and affine maps,

centered reconstruction, rank sensitivity, principal angles, and a separately identified full-domain oracle. Centering by training means prevents a large constant feature component from dominating reconstruction error.

Table 23: Centered held-out reconstruction error (%) for a train-fitted map from later to earlier residuals. Healthy controls match interval length. Reference intervals span 6–508 updates and occur at different optimization stages, so comparisons are within each seed.

Seed	Reference to captured failure	Healthy interval	Adjacent failure update
0	43.63	32.66	1.88
1	43.05	10.82	4.63
2	9.05	2.76	1.10
3	45.73	14.30	10.66
4	25.60	1.18	25.63

Positive controls plant exact residual gauges into the trained network. Eight controls cover condition numbers 1, 3, 10, and 100 in float32 and float64; twenty-four further controls plant observed fitted matrices with their measured orientation and conditioning. Exact compensation and train-fitted compensation recover 100% in every fixed-condition control. For the three fresh adjacent transitions, observed forward-map errors are 4.65%, 1.10%, and 10.77%, against 0.0004–0.0006% for their planted controls. These errors reject exact global equivalence for the measured pairs. Healthy evolution also departs from exact equivalence (Table 23).

Approximate transport remains useful: inverse-linear compensation reaches 95.13%, 100%, and 94.37% at the three fresh first events. Reference timing matters. Seed 0’s head from step 17,200 scores 14.05% on the collapsed residual, while the head from 17,493 scores 96.39%. Its geometric discrepancy from the distant reference is already present before the final failing update. Moreover, the complete real task-only Fourier coefficient matrices have shape 113×128 ; when both have full row rank, an invertible map between them always exists. For the seed-4 generalization-drop pair at steps 16,000 and 17,000, a fitted task-only map aligns that projection exactly while scoring 4.62% on raw held-out features. This pair is separate from the adjacent failures above.

One-time repairs also have limited persistence. Eighteen continuations resume two captured failures for 500 updates under nine head/state interventions each. Every successful repair retaining optimizer state falls below 90% training accuracy on the next update. Clearing the head’s Adam state improves eventual recovery but permits the immediate relapse. These tests separate instantaneous decodability from stability under continued optimization.

J.3 ARITHMETIC INTERVENTIONS

The accurate loss evaluates $\log(1 + \sum_{j \neq y} \exp(z_j - z_y))$ with $\log 1p$ whenever the correct-class logit is maximal; other examples use ordinary cross-entropy. Model storage, backward matrix arithmetic, and Muon’s update remain float32. This avoids cancellation in the small loss and correct-class derivative, following the numerical concerns of Prieto et al. (2025). All five step-6,000 sources already satisfy the six-check grokking criterion and already contain measurable gradient error. The intervention acts on an existing post-grokking trajectory.

The specificity study branches each seed at step 15,000 and runs to 20,000 with three corrections. Accurate CE gives 0/5 joint failures; target-derivative repair gives 0/5; per-example zero-sum projection gives 0/5. Every original trajectory has a captured failure in this window. All fifteen corrected branches finish at 100% test accuracy; the minimum measured test accuracy across these branches is 99.978%.

For stock logit gradient g , target repair retains each wrong-class entry and sets $g_y = -\sum_{j \neq y} g_j$, leaving the reported forward loss unchanged. Projection replaces g by $g - \mathbf{1}(\mathbf{1}^\top g)/p$. Both corrections compute their class sums in float64 and cast the result back to float32. Derivative errors are relative L_2 errors against the accurate float64 derivative evaluated on the identical float32 logits. The readout-gradient reference casts this logit derivative to float32 before backpropagating through the shared float32 network. At the fixed seed-4, step-16,000 audit, target repair reduces relative deriva-

Table 24: Matched arithmetic branches from step 6,000 through 30,000. Original columns show a captured joint failure. Accurate-CE branches have no joint failure; all finish at 100% train and test accuracy. Minima include the starting state.

Seed	Original event step	Original event test (%)	Accurate minimum test (%)
0	17,494	31.872	99.754
1	17,721	9.095	99.843
2	18,069	0.929	98.859
3	16,308	53.597	99.978
4	16,056	66.327	99.843

tive error from 32.02% to 4.19×10^{-7} , whereas projection leaves 31.88%. At step 20,000, projected derivatives retain 52.94–86.07% relative error despite a class-sum residual below 2.1×10^{-8} relative to the reference norm. Thus restoring the class-sum identity can prevent the observed failures over this window even with substantial remaining derivative error. Numerical feature inflation and this projection control build on Liu et al. (2026).

At step 15,000, stock logit-gradient errors are 10.56–21.86% and readout-gradient errors 18.95–43.42%. Passing a class-common gradient error through the same saved Adam state produces a class-centered update difference: its centered norm fraction is 45.85–73.37%, versus $(2.8\text{--}4.1) \times 10^{-16}$ for scalar SGD, at float64 roundoff. These descriptive norm ratios are nonadditive. Accurate-CE branches also have smaller feature means: original mean norms at step 15,000 are 2.22–3.83 times their matched corrected values.

Persistence through the original training horizon. We resume each corrected step-30,000 state through step 100,000, matching the original depth-1 horizon. All model parameters, optimizer buffers, and the PyTorch random state are restored. Twelve direct updates agree bitwise with both the monitored runner and a six-plus-six resumed continuation for all six sources. Each extension adds 70,000 updates, with training accuracy recorded at every state and test accuracy every 100 steps and whenever training accuracy is below 90%. The final and minimum-accuracy checkpoints reproduce their logged metrics exactly.

Table 25: Corrected-loss runs through step 100,000. Accuracies are percentages. Minima combine the earlier prefix with the extension, counting the shared step-30,000 state once. Addition is corrected from step 6,000; subtraction uses corrected CE from initialization and its minima start at grokking confirmation.

Run	Monitoring start	Min. test	Final test	Joint event
Addition, seed 0	6,000	99.754	100.000	none
Addition, seed 1	6,000	99.843	100.000	none
Addition, seed 2	6,000	98.859	100.000	none
Addition, seed 3	6,000	99.978	100.000	none
Addition, seed 4	6,000	99.843	100.000	none
Subtraction, seed 0	10,500	96.879	99.955	none

All six extensions complete without a joint event. Thus correction suppresses the observed failures through the original training budget. This addresses postponement beyond step 30,000 within that budget; behavior after 100,000 and brief test-only excursions between scheduled evaluations remain outside the measurement. The original stock trajectories already have captured failures by 30,000, which establishes their positive binary incidence at the longer horizon. Their historical monitoring supports no comparison of event counts or durations with these dense branches.

J.4 MEAN AMPLIFICATION AND FINITE-STEP OVERSHOOT

At the checkpoints immediately before the five original failures, feature means over the full operand grid have norms 10,374–12,203. The main-text intervention in Eq. 1 uses the post-update representation’s training mean and adds each component to the rescued logits $W_{U,0}h_1$. Applying the

training-estimated mean unchanged to test examples, the mean term alone gives 32.11%, 9.42%, 0.93%, 53.34%, and 66.07% accuracy, close to the full failing updates in Table 24. The centered term alone retains 95.26%, 92.48%, 100%, 87.20%, and 98.46%. Relative reconstruction error is below 3×10^{-16} . This evaluation-only intervention establishes local sufficiency of the product of feature mean and readout displacement. Feature growth and optimizer history remain coupled before that step.

Replacing gradients only at the failing update with accurate gradients preserves all five failures. The state accumulated beforehand matters. For each actual complete parameter displacement $\Delta\theta$, the accurate training gradient satisfies $\nabla L^\top \Delta\theta < 0$, including each group’s contribution. Scaling the displacement by 0.001 or 0.01 lowers training loss in every seed; taking the full step raises it sharply. Both interpolation endpoints reproduce saved parameters exactly. The captured updates locally descend and overshoot at finite scale.

J.5 SUBTRACTION AND NORMALIZED ARCHITECTURE

Four additional runs use seed 0 and a fixed 30,000-step horizon: a stock/accurate-CE pair on subtraction and a pair on addition with gain-free RMS normalization before attention, before the MLP, and before the readout ($\epsilon = 10^{-6}$). Parameter initialization and optimizer settings are matched, with no added parameters or hyperparameter search.

Table 26: Operation and architecture controls. Confirmation is the end of the six-check grokking streak. First-event test accuracies and post-grokking minima are separate measurements; both RMS runs subsequently recover to 100%. Each row is one seed.

Condition	CE	Confirmation	First event	Event test (%)	Min. test (%)
Subtraction	Stock	9,300	15,281	5.851	5.851
Subtraction	Accurate	10,500	none	n/a	96.879
RMS addition	Stock	15,700	18,515	85.658	0.481
RMS addition	Accurate	19,900	28,495	0.884	0.526

Subtraction reproduces acute readout sensitivity: updated features with the preceding head score 98.81%, the mean term alone 5.84%, and the centered term 98.76%. At step 30,000, accurate CE has no joint event and test accuracy is 99.966%. Its continuation through 100,000 is reported in Table 25.

In the RMS runs, the embedding update causes the acute loss of accuracy. In the accurate-CE event, test accuracy falls from 96.208% to 0.884%; its logit derivative has relative error 9.48×10^{-8} and no vanished correct-class entries. Updating only the head retains 96.208%, updating only embeddings gives 0.940%, and freezing the head still gives 0.895%. The minimum reaches 0.526% at step 28,503; the stock-RMS minimum is 0.481% at 28,540. Exact update replay passes. Both RMS event directions locally descend and lower loss at scales 0.001 and 0.01, then overshoot at full scale. This counterexample bounds the readout/feature-mean account to the measured unnormalized cases and rules out universal prevention by accurate loss arithmetic. The completed horizons support those finite claims; longer runs and population frequencies remain outside their scope.




Figure 8: One-step parameter interventions at the first detected joint failure in each operation or architecture control, all at seed 0. Each bar uses the preceding checkpoint with only the indicated parameter group replaced by its following value; “Full” applies all three groups. In subtraction, the head update reproduces most of the damage. In both RMS cases, the embedding update is destructive while the head update preserves accuracy. The accurate-CE RMS failure occurs with an accurate loss derivative. These are first-event measurements at the displayed steps; the later post-grokking minima are reported separately in Table 26.

K TASK STRUCTURE IN FOURIER FILTERING

K.1 FULL-FAMILY PROJECTION IS AN ORBIT AVERAGE

Write a Fourier mode as $\chi_{k,l}(a, b) = \exp(2\pi i(ka + lb)/p)$. Averaging its translations gives

$$\frac{1}{p} \sum_{t=0}^{p-1} \chi_{k,l}(a + t, b - t) = \chi_{k,l}(a, b) \frac{1}{p} \sum_{t=0}^{p-1} \exp\left(\frac{2\pi i(k-l)t}{p}\right) = \mathbf{1}_{k=l} \chi_{k,l}(a, b). \quad (3)$$

Linearity proves Eq. 2 for each feature channel. The surviving modes include $(0, 0)$. For subtraction, averaging $(a + t, b + t)$ instead retains $k + l = 0$ modulo p . With a linear head, $W_U Ph = P(W_U h)$; an affine head also commutes with this average because averaging preserves its constant bias. Excluding the constant mode gives the orbit average minus the global mean.

The full-family filter therefore supplies the equivalence relation defining the answer. It can pool logits from training and held-out examples with that answer. This is part of the intervention itself, even though no labels are passed to the FFT routine. Its accuracy measures an answer-conditioned average and requires separate evidence to identify the algorithm implemented by the unfiltered model.

K.2 A LOOKUP-TABLE MEMORIZER

We use the repository’s five seeded splits, each with 3,830 training and 8,939 held-out examples at $p = 113$. Let T denote the training set and e_s the one-hot vector for class s . The predictor stores $z(a, b) = e_{a+b}$ on T and outputs zero on every unseen pair. If n_s is the number of training examples with sum s ,

$$(P_{\text{add}} z)(a, b) = \frac{n_{a+b}}{p} e_{a+b}. \quad (4)$$

Every answer orbit is covered in each split, so filtering yields 100% accuracy. For a uniform split of $m = 3,830$ out of $N = p^2$ pairs, an orbit is uncovered with probability $q = \binom{N-p}{m} / \binom{N}{m}$. A union bound gives $\Pr(\text{some orbit uncovered}) \leq pq < 2.89 \times 10^{-16}$. Before filtering, all held-out logits are tied. We use the smallest class index as the deterministic tie rule, matching the implementation’s argmax; uniform random tie breaking has expected accuracy $1/113 = 0.885\%$.

Table 27: Addition memorizer on the original seeded splits. Correct counts use the 8,939 held-out examples. The filtered logits have a unique correct maximum on every example.

Seed	Raw correct	Raw accuracy (%)	Filtered correct	Filtered accuracy (%)
0	79	0.884	8,939	100.000
1	68	0.761	8,939	100.000
2	79	0.884	8,939	100.000
3	81	0.906	8,939	100.000
4	77	0.861	8,939	100.000

The same control on subtraction reaches 100% after the corresponding filter in all five seeds. Omitting the constant mode also retains perfect accuracy in these controls. A seeded bijection of output labels preserves the perfect filtered result when evaluated against those relabeled classes, because the projection still pools within answer orbits. Averaging training donors alone is sufficient; averaging only held-out donors gives zero logits for the memorizer.

These controls establish that full-family filtered accuracy can be perfect without a learned generalizing arithmetic algorithm. The 95.14% filtered accuracy at AdamW step 3,000 therefore requires independent evidence to establish when an internal arithmetic algorithm was learned. Evidence for surviving information at the collapsed checkpoints comes from the raw-feature decoder tests and adjacent-state swaps in §4.2.

K.3 SPARSE SUFFICIENCY, ABLATION, AND PERTURBATION CONTROLS

A second lookup stores $(p/n_c)e_c$ for a training example with label c , using training-label counts only, and zero for unseen pairs. Its full projection is exactly e_{a+b} . Retaining any one nonzero diagonal




Figure 9: Task structure supplied by Fourier filtering at $p = 113$. **(a)** The lookup memorizes the 3,830 training answers and returns zero on the 8,939 held-out inputs. Full-family filtering with DC gives 100% accuracy. Family ablation retains DC and yields 0% held-out accuracy and 29.99% full-grid accuracy. Bars show means over five splits; whiskers span their range. **(b)** Training-count normalization allows any one diagonal conjugate pair plus DC to classify every input correctly. The ordinary lookup’s pair powers are mathematically tied; its prefixes use increasing frequency as a fixed tie break. Lines show means and shading spans the seed range.

conjugate pair and DC gives class- c logits

$$z_k(s, c) = \frac{1 + 2 \cos(2\pi k(s - c)/p)}{p}, \quad s = (a + b) \bmod p. \quad (5)$$

For prime $p = 113$ and nonzero k , the unique maximum is at $c = s$. All 56 individual pairs give 100% training, held-out, and full-grid accuracy on all five splits (280 checks against the closed form). Thus sparse filtered sufficiency also admits a lookup-only predictor. The construction can use a 113-dimensional lookup representation with an identity head, padded to the transformer’s 128-dimensional residual width.

The ordinary one-hot lookup has equal mathematical power in all diagonal conjugate pairs. With a fixed ascending-frequency tie break, five-pair filtering gives 23.82–28.57% held-out accuracy; sorting the numerically computed powers instead gives 63.69–87.46%. These ranges expose the sensitivity of the selected frequencies to floating-point tie breaking in this control. Figure 9 uses the fixed ordering. The count-normalized lookup reaches 100% at every tested prefix size 1, 2, 3, 5, 10, 20, 56.

The original family-ablation convention is $z - Pz + \bar{z}$, retaining DC. For both lookups, it preserves all training predictions and gives 0% held-out accuracy, hence $3,830/12,769 = 29.9945\%$ full-grid accuracy. The trained models’ near-chance full-grid ablations in Table 4 contain additional information absent from this simple control. Full-family sufficiency and low held-out ablation accuracy can nevertheless occur together under memorization.

We also apply 20 fixed-seed replicates per split and lookup for pair-global phases, channelwise phases, and derangements of the 56 pair coefficients, preserving DC, conjugate symmetry, and total nonconstant family power. For the ordinary lookup, mean held-out accuracies of the isolated family are 1.0755%, 0.9426%, and 3.0680%, respectively; for the count-normalized lookup they are 0%, 0.8850%, and 3.1555%. Each mean pools 20 perturbations per split across five splits; raw replicate distributions are released. The sensitivity observed in the trained models remains a functional measurement of their projected components. These controls show why that sensitivity alone leaves the internal algorithm unidentified.

K.4 RESTRICTING THE AVERAGING DONORS AT CAPTURED FAILURES

Using both endpoints of the five adjacent failure pairs in Table 22, we recompute logits from the saved raw residuals and native heads, multiplying in float64. We average only held-out inputs with the same answer, excluding the evaluated input itself. This removes both training donors and self contribution. Input order, labels, and masks are checked against the original seeded splits.

Table 28: Held-out accuracy (%) before and after each captured update. “Average” uses held-out donors with the same answer and excludes the evaluated example. The operation still supplies the grouping.

Seed	Before		After	
	Raw	Average	Raw	Average
0	95.49	100.00	31.87	30.46
1	94.14	100.00	9.09	0.93
2	99.99	100.00	0.93	0.88
3	89.54	99.12	53.60	65.29
4	91.34	98.18	66.33	71.25

Before these updates, high accuracy survives removal of training donors, establishing task-aligned signal in the logits of unseen examples beyond the ordinary lookup control. The donor-restricted operation still uses the answer orbits and remains a task-informed diagnostic. These checkpoints are distinct from the early AdamW checkpoint at step 3,000; this analysis makes no claim about its donor dependence.

Independent FFT, translated-grid sums, and orbit reductions agree in float64 at moduli 7 and 113. We also check projector idempotence, constant-mode conventions, readout commutation, coverage of every answer orbit, and the single-pair closed form. All scripts, protocols, integer counts, perturbation outcomes, and input hashes accompany the release.